

# (python) Programming in APC

---

## Lecture 1 : Computer architecture and the Cluster

Tommaso Ronconi  
AA 2023-24



# Outline for today

---

## Why clusters and why should you care?

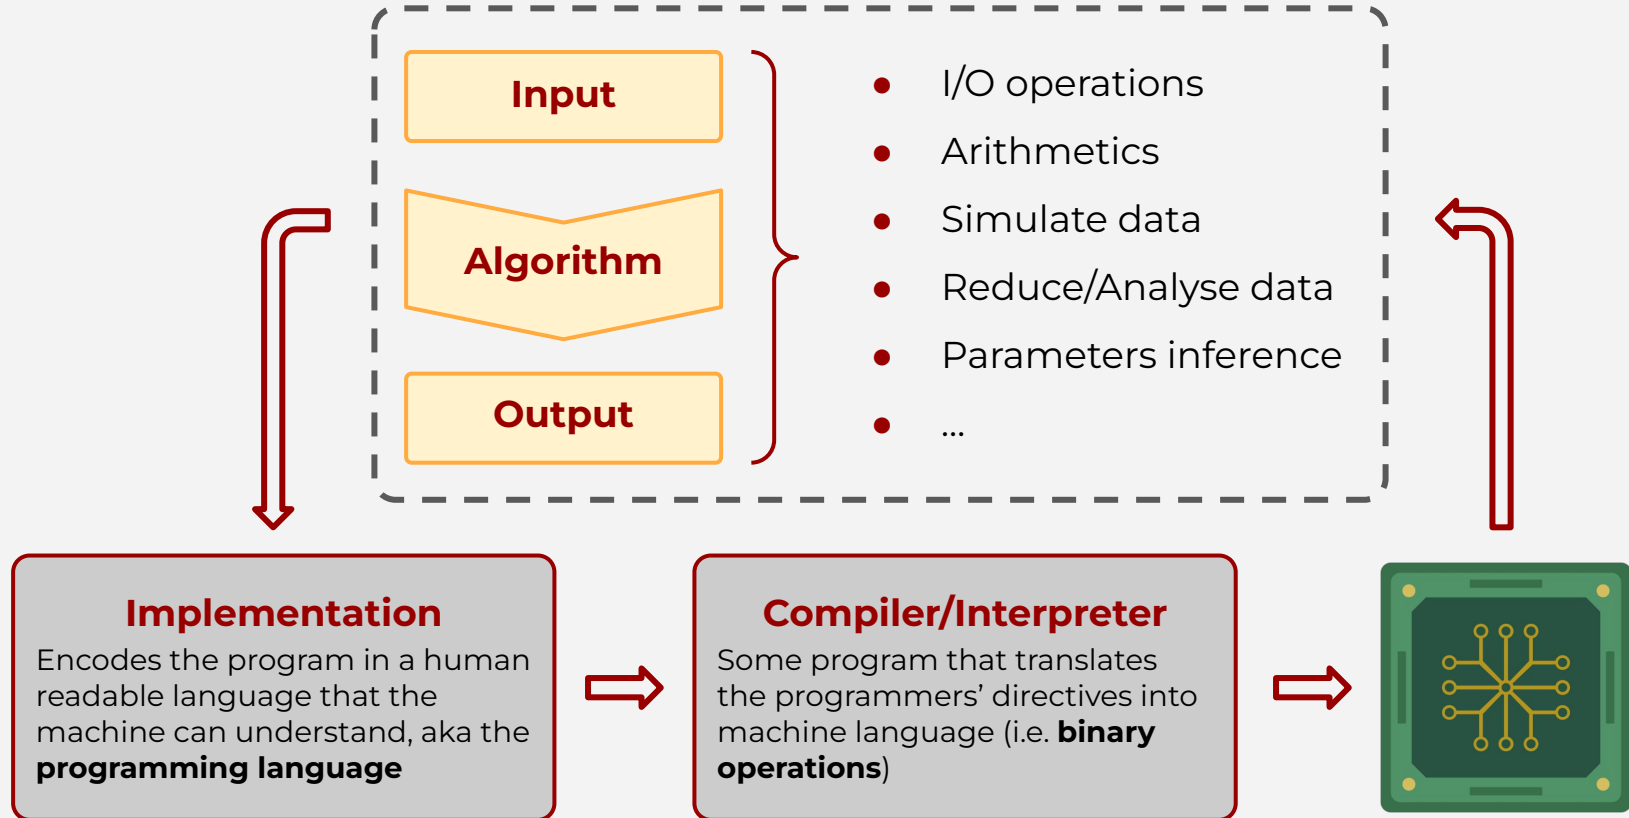
- Familiarize with what a computer **program/application** does
- Primer to **processors** architecture and to **cluster architecture**
  - Thinking in **parallel**
  - **Communicate** with the cluster + **etiquette**

```
$ OMP_NUM_THREADS=N mpirun -np M program.x
```

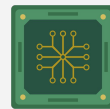
---

# What is a program?

... or an application or a script, whatever



# “There ain’t no such thing as a free lunch”



**Before: Out-of-the-box** increase in performance, new hardware == more power

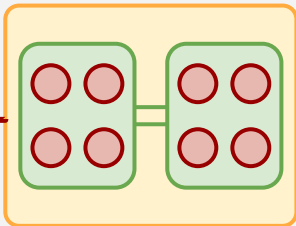
**Now:** Increase of performance comes at the cost of implementing with **hardware awareness**

**Physical and Technological limits**  
to computational power



**PARALLEL  
ARCHITECTURE**

User



Mother-board



Socket = CPU = Processor



Core (sometimes = “processor” -.-)

**Each** element provides different levels of parallelism



# How to think in parallel?

---



## DOMAIN DECOMPOSITION

- **crunch more numbers** (more cores on the same problem)
- **distribute the domain** (more RAM on the same problem)

## TASK DECOMPOSITION

- **reduce starvation** (elements of the machine not executing instructions)
- **reduce overhead** (distribute the parts of the path-to-solution that can be executed out of order)

# Levels of parallelism



- **serial execution** but actually, **superscalar**:

- instruction-level parallelism (ILP)

**c** = **a** \* **b**

**f** = **d** + **e**

- hyper-threading (next slide)

- pipelines:

fetch instruction > load data > ~~execute~~ computation > store data

- vectorization

for i in indices/3:

c[i] = a[i] + b[i]

c[i+1] = a[i+1] +

b[i+1]

c[i+2] = a[i+2] +

b[i+2], computation > store

- **multi-threading/processing**: shared memory archs. open/close parallel regions from within a serial application

- **distributed computing**: parallel copies of the program run independently and communicate through a **message passing interface (MPI)**

- **accelerators**: GPU/TPU (hardware designed for matrix operations, not covered in this lectures)

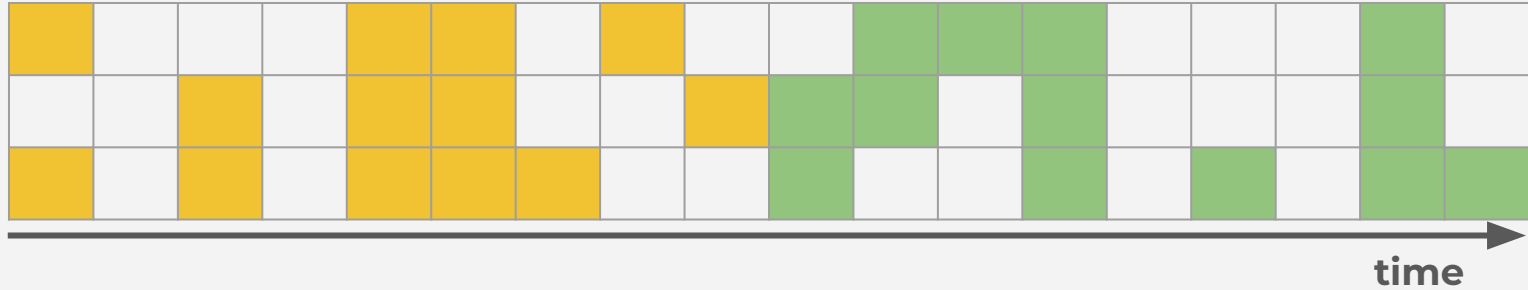


# Levels of parallelism: hyper-threading



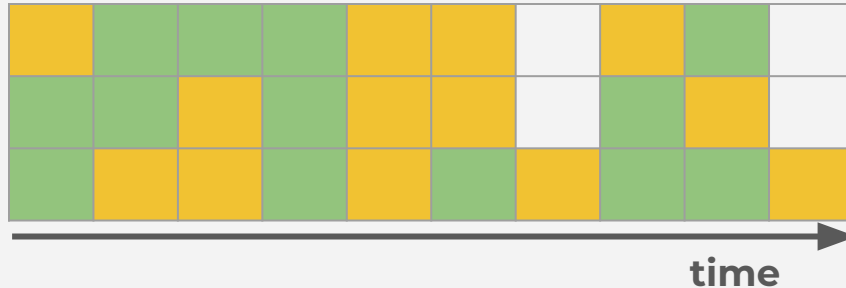
## Standard Core:

One thread after the other



## Hyper-threading Core:

The two threads run together

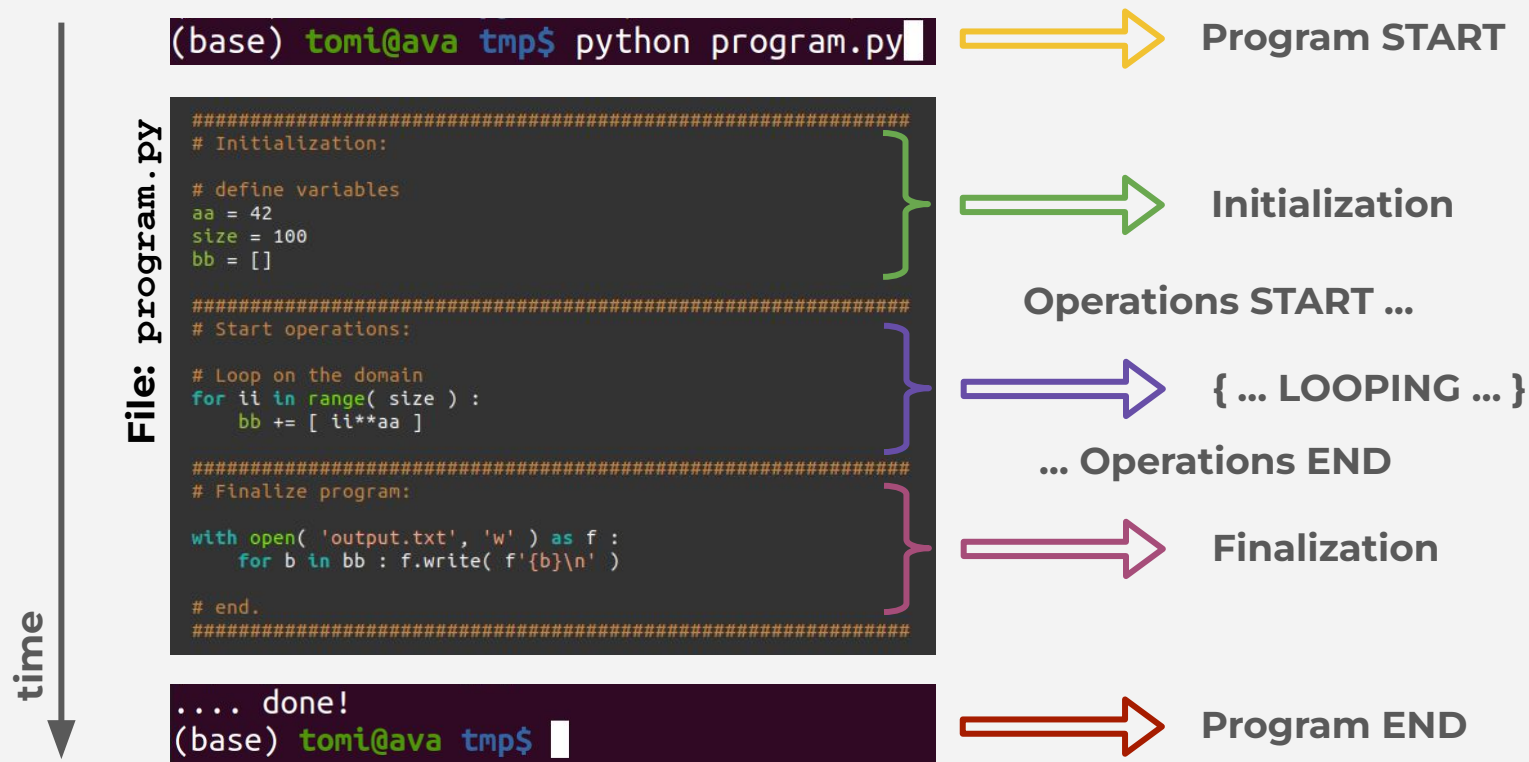
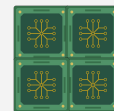


Idle

Thread 1

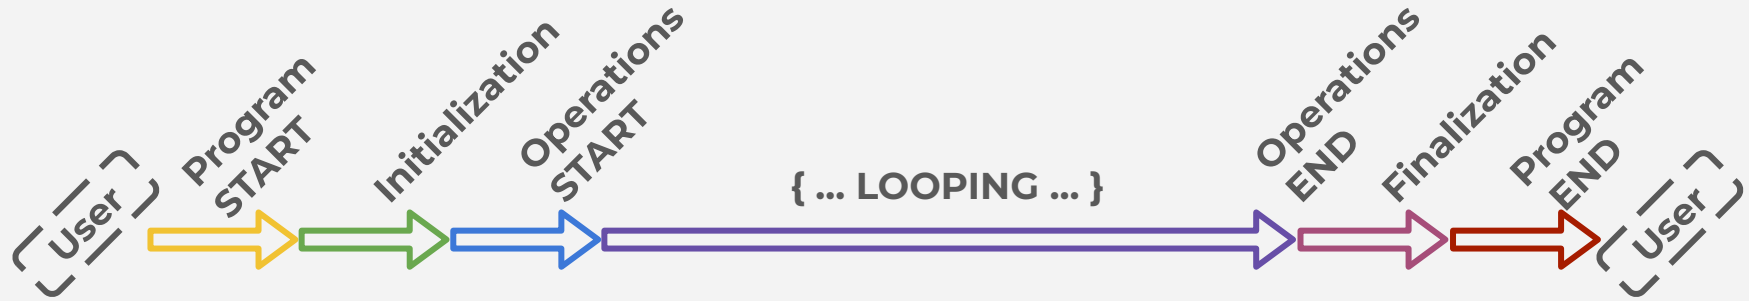
Thread 2

# Levels of parallelism: a toy model

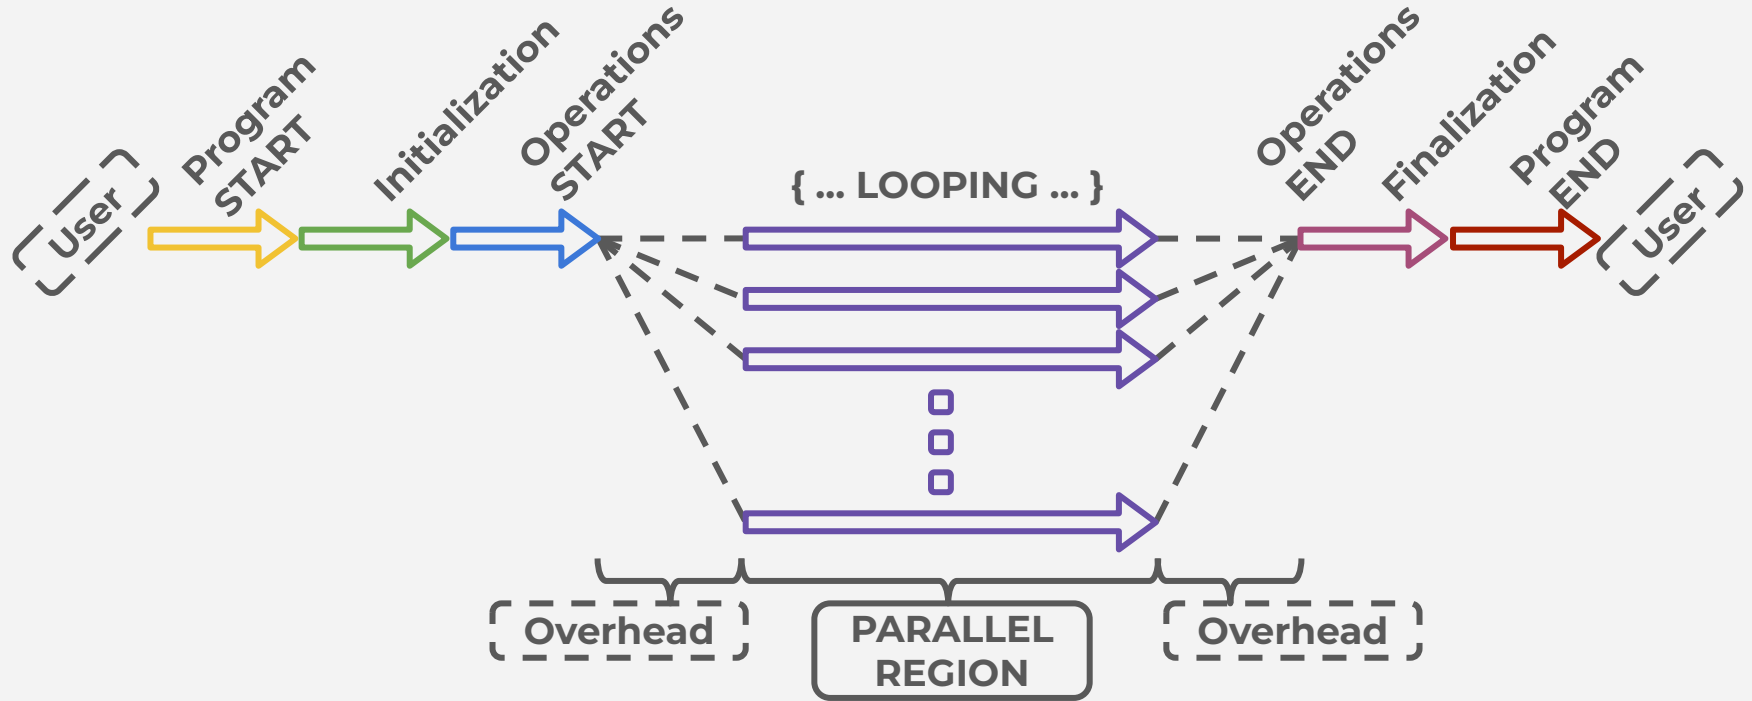


# Levels of parallelism: serial execution

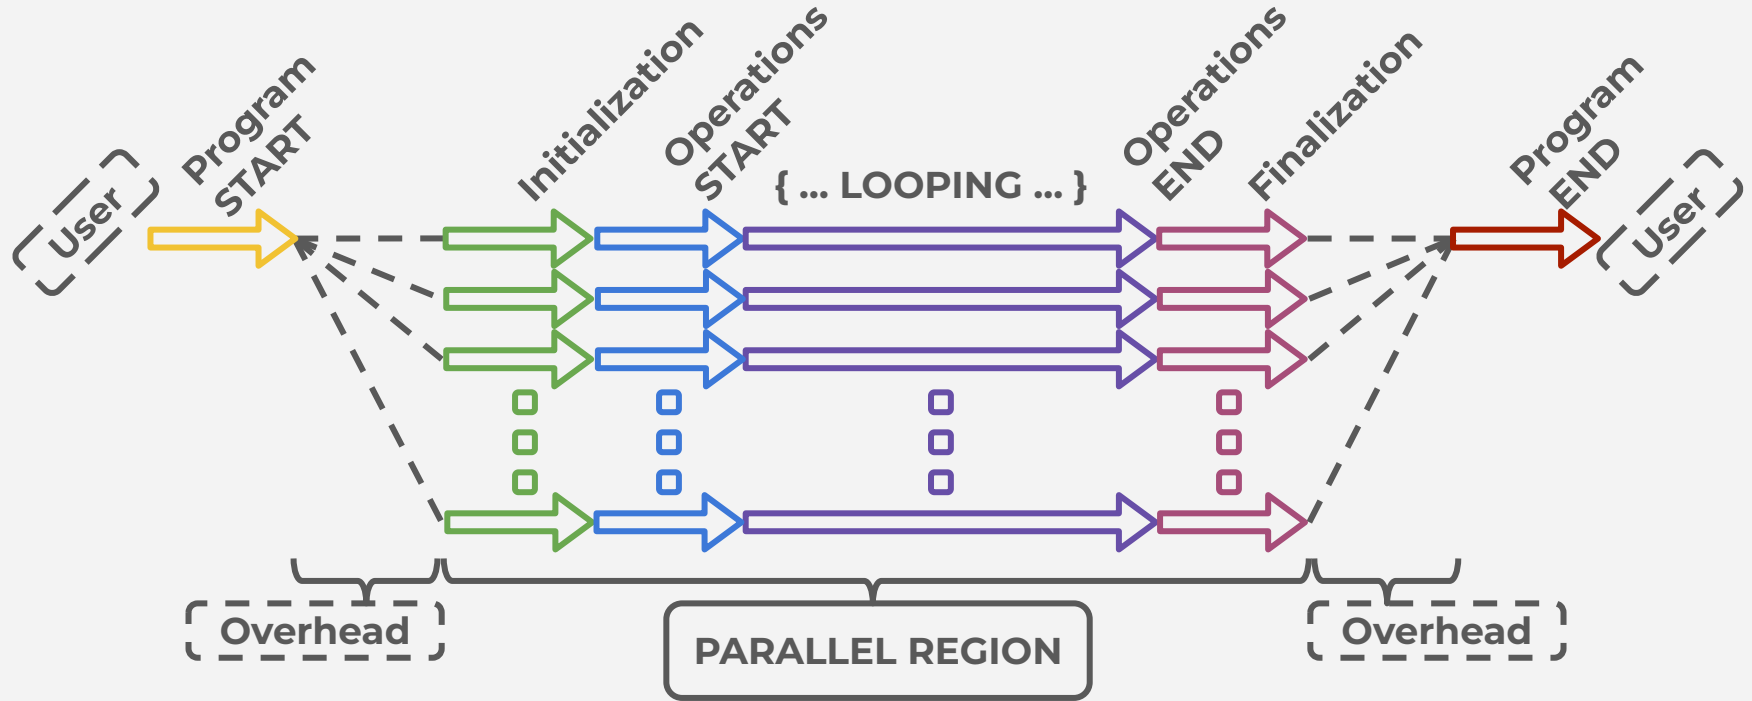
---



# Levels of parallelism: multi-processing



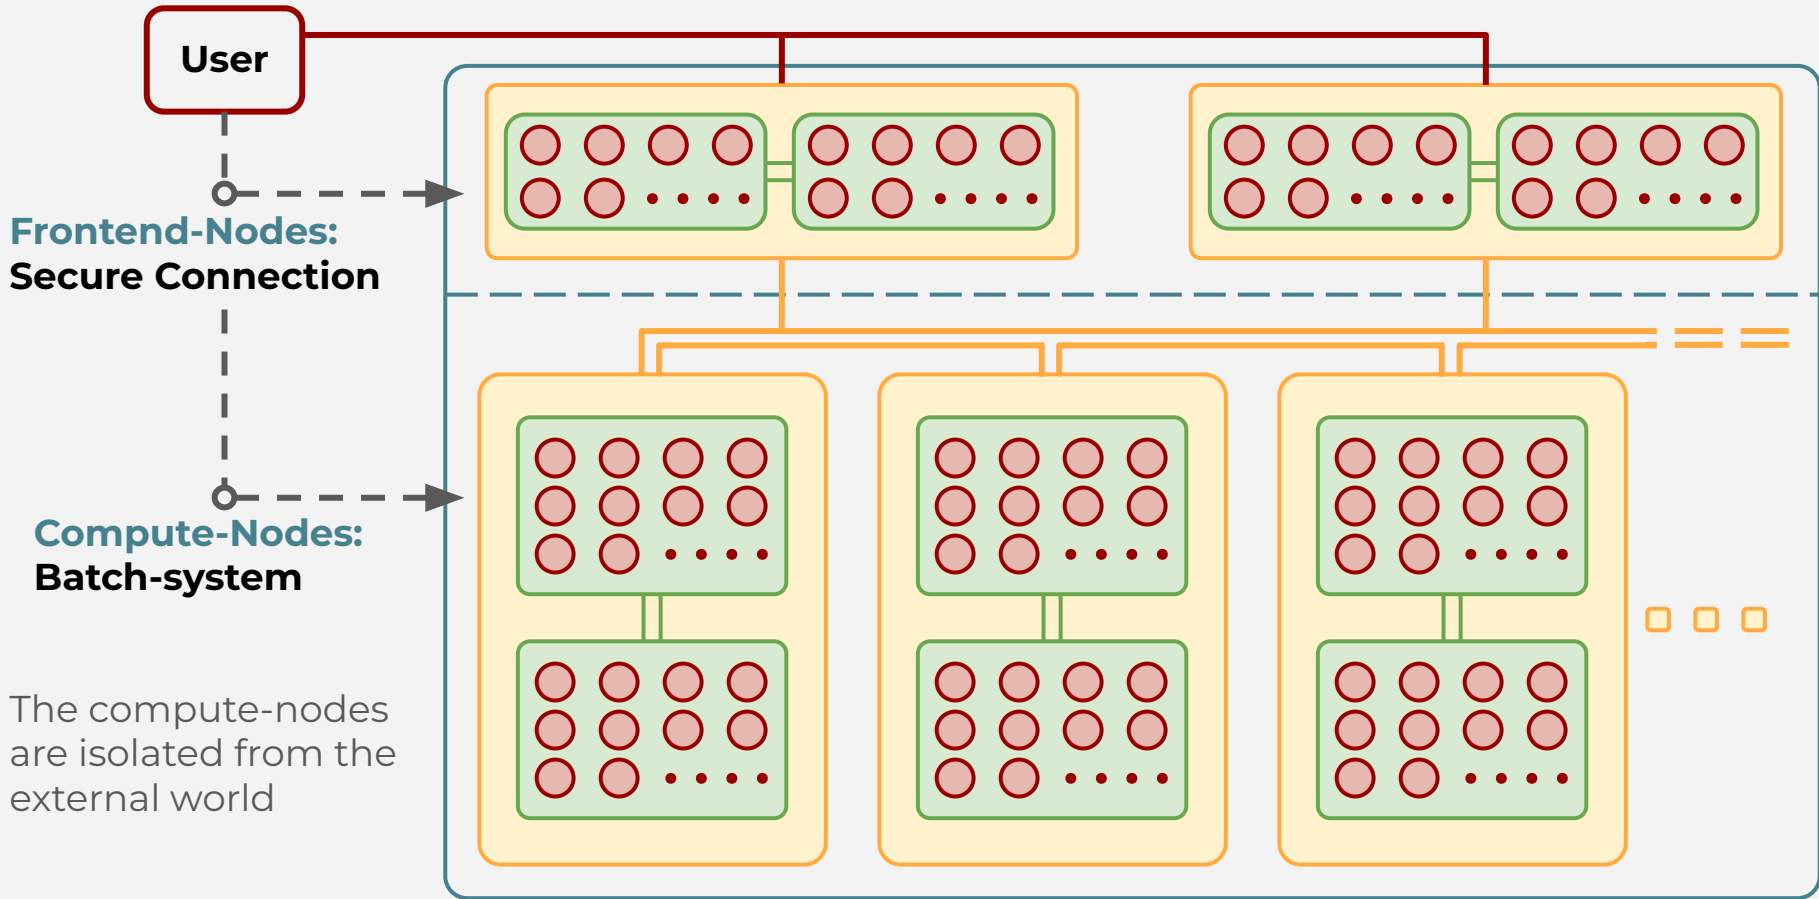
# Levels of parallelism: message passing



ssh user@powerfull\_machine:/CLUSTER\_BASICS

---

# What is a Super-computer today?



# The Secure-Shell protocol (SSH)

---



Provides **File access/transfer/management** through a double-ended exchange of secure-keys

## Connect to the Remote-Server:

```
$ ssh user@server_domain_name (or IP_address)
```

## Disconnect:

```
$ logout
```

## Copy-stuff from/to Remote-Server:

```
$ scp user@server:/remote/path/to/file.ext /local/path/to/dir/ (from remote)
```

```
$ scp /local/path/to/file.ext user@server:/remote/path/to/dir/ (to remote)
```

## Also works:

```
$ rsync -av [position_1] [position_2]
```

(Synchronization is better when you have large files)

The **Secure-Shell** is directly accessible only from the **front-end nodes**  
Communication with the **compute-nodes** is performed through the **Batch-System**.



# The Scheduler/Batch-System



if people understood socialism a batch system wouldn't be necessary

Provides **User access to compute facility** attempts to provide a **fair-share of resources**

**Flavours:** **Torque** (old-school) or **Slurm** (new-school)

... the two on the surface differ for the commands

Imagine a dystopian future where the world resources are managed by an AI deciding who is getting what and when.

**To get resources from a cluster you need to ask the batch-system**

TIME

MEMORY

POWER

# Interactive jobs

You can ask the batch system to open a shell session on some compute node:

```
$ srun --partition=regular1 --time=00:30:00 --nodes=1 --mem-per-node=40gb --pty bash -i
```

The diagram shows the command `$ srun --partition=regular1 --time=00:30:00 --nodes=1 --mem-per-node=40gb --pty bash -i` with brackets underneath each part, each pointing to a label:

- `srun`: Ask slurm resources
- `--partition=regular1`: where?
- `--time=00:30:00`: for how long?
- `--nodes=1`: how much power?
- `--mem-per-node=40gb`: how much memory?
- `--pty bash -i`: Make it interactive

# Submit jobs

Or you can ask to put the job in the queue, when your turn arrives execution is automated

The diagram shows the command `#SBATCH --partition=regular1` with an arrow pointing to **where?**, `#SBATCH --nodes=1` and `#SBATCH --ntasks-per-node=20` with a bracket pointing to **how much power?**, `#SBATCH --ncpus-per-task=1` with an arrow pointing to **how much power?**, `#SBATCH --mem-per-cpu=4096` with an arrow pointing to **how much memory?**, `#SBATCH --time=00:30:00` with an arrow pointing to **for how long?**, and `/path/to/the/program_to_run.ext` with an arrow pointing to **What has to be done.**

```
#!/bin/bash

#SBATCH --partition=regular1
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=20
#SBATCH --ncpus-per-task=1
#SBATCH --mem-per-cpu=4096
#SBATCH --time=00:30:00

/path/to/the/program_to_run.ext
```

**Launch the slurm-job:**

```
$ sbatch script_name.sh
```

**Show job status:**

```
$ squeue -u user_name
```

**Cancel job:**

```
$ scancel job_id
```

# Ulysses: the SISSA cluster

---



I could have prepared a slide but sharing this link is simpler for all of us.

<https://ulysses.readthedocs.io/index.html>

So long, and thanks for all the fish

---